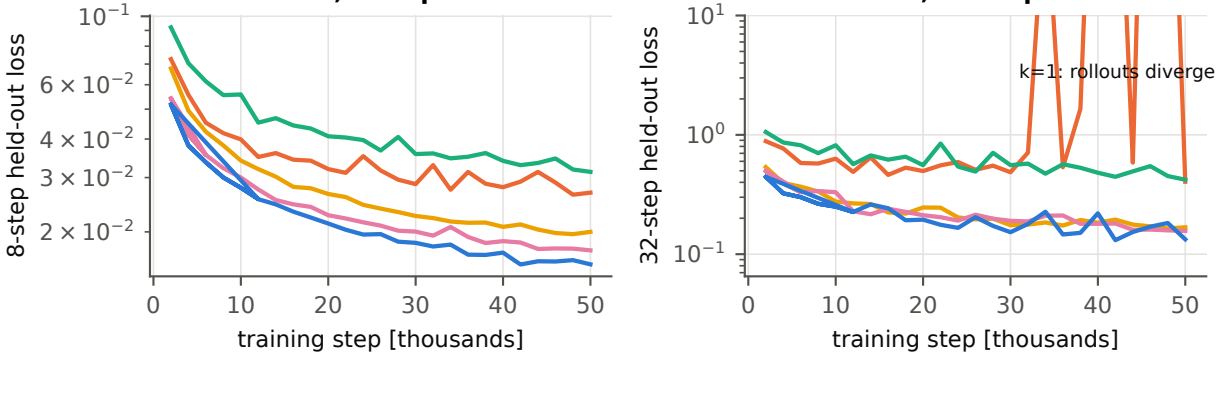
Validation, 8-step windows
8-step held-out loss
10^{-1}
6×10^{-2}
4×10^{-2}
3×10^{-2}
2×10^{-2}
0
10
20
30
40
50
training step [thousands]
Validation, 32-step windows
32-step held-out loss
10^{1}
10^{0}
10^{-1}
k=1: rollouts diverge
0
10
20
30
40
50
training step [thousands]
k=1 training
absolute targets
width 1024
unseen policy
main: k=8, 2048×3, Δ